


Figure 5: Multi-Agent Scaling Study (100 memories/agent, no-LLM mode)

(a) Query Latency vs. Agent Count



(b) Memory Footprint & Scaling

